

Lab 1: Introduction to MPI

EEL 6763 – Spring 2023

Part 1: Matrix Multiply

Run the given mat_mult.c as is (without changing anything) **with 8 ranks** to obtain the results (as a baseline).

Then modify the mat_mult.c to do the following; and again run it **using 8 ranks**:

Each worker rank will:

- Determine two times: (using the now() and tdiff() functions)
 - wcalcTime: the actual time spent on performing its calculation of matrix C.
- Send its wcalcTime to the master rank.

The master rank will:

- Determine time_spent: (time from sending the data to workers until the master have all the results). By the way, this is already in the original mat_mult.c.
- Print (at the end):
 - wcalcTime for each worker rank
 - maximum wcalcTime
 - sum of the wcalcTime for all worker ranks
 - time_spent (in the master rank, again already in the original code)

Part 2: Monte-Carlo Integration and MPI

One method of numerically estimating integrals is by using Monte-Carlo simulation. Consider the following integral $g(a,b)$ and its estimate $h(x)$:

$$g(a,b) = \int_a^b f(x) dx \qquad h(x) = \frac{(b-a)}{N} \sum_{i=1}^N f(x_i)$$

The numerical solution to $g(a,b)$ can be estimated using a uniform random variable x that is evenly distributed over $[a, b]$ (i.e., between a and b). The estimate $h(x)$ will converge to the correct solution as the number of samples N grows. Since each sample is independent, the calculation can be easily parallelized (embarrassingly parallel). Using this method, write a short MPI program that will use N samples to calculate:

$$\int_a^b \frac{8\sqrt{2\pi}}{e^{(2x)^2}} dx$$

For this problem you must provide code for three functions, `init_rand_seed()`, `estimate_g(...)` and `collect_results(...)`, which will be used by the following `main(...)` function:

Important note: Use the main function as it is defined below. It cannot be changed (e.g., you should not be doing any calculation, send/recv, or printing in the main function).

```
int main(int argc, char **argv)
{
    double result = 0.0;
    MPI_Init(&argc, &argv);
```

Lab 1: Introduction to MPI

EEL 6763 – Spring 2023

```
float lower_bound = atof(argv[1]);
float upper_bound = atof(argv[2]);
long long int N = atof(argv[3]);

init_rand_seed(); // using srand()
result = estimate_g(lower_bound, upper_bound, N);
collect_results(&result);

MPI_Finalize();
return 0;
}
```

The following function prototypes must be used for your functions:

- `double estimate_g(double lower_bound, double upper_bound, long long int N);`
- `void collect_results(double *result);`

Other specifications:

- The total number of samples to calculate **N** (to be split among all MPI nodes), as well as the bounds of the integral **a** and **b**, will be provided through command-line arguments.
- Every MPI rank will generate its own random numbers. Ensure that each rank uses a different starting seed for its random number generator - this should be the only thing that occurs in , `init_rand_seed()`.

Then in the `estimate_g` function, you should use the `rand` function to generate the random numbers. Hint: While debugging print out the random numbers to verify that they are indeed random.

- Each rank should return a single value, which should be combined (i.e., sum) in the master rank in order to compute the final integral.
- Use only the following functions:
 - `MPI_Init`
 - `MPI_Comm_rank`
 - `MPI_Comm_Size`
 - `MPI_Send`
 - `MPI_Recv`
 - `MPI_Finalize`

Part 3: MPI Reduce

Rewrite the previous function (from Part 2) to use `MPI_Reduce` instead of `MPI_Send` and `MPI_Recv`.

Lab 1: Introduction to MPI

EEL 6763 – Spring 2023

SUBMISSION INSTRUCTIONS

- Make sure your name is at the top of every file.

You are to submit 2 files on Canvas:

- (1) Create a directory named **Lab1**. Use the following structure and zip the entire directory and submit the zip file on Canvas using the following name: lastNamefirstNameLab1.zip

Lab1/

partA/

.c file

batch script

output file

partB/

.c file

batch script

7 output files: a = -10 ; b = 10 (keep a and b constant);

- Keep N = 100000, vary R (number of ranks) = 4, 8, 16, 32
- Keep R = 32, vary N= 100, 1000, (10000), 100000
- Name each output file appropriately: e.g., SendRecvR32N100, SendRecvR8N100000)
- Also, to maximize your credit (especially partial credit), use printf statements like they were used in the mat_mult code (e.g., number of elements per rank, partial results sent back by each rank, etc.)

partC/

.c file

batch script

7 output files: same specifications as Part B.

- Name each output file appropriately: e.g., ReduceR32N100, ReduceR8N100000).
- Also, to maximize your credit (especially partial credit), use printf statements like they were used in the mat_mult code (e.g., number of elements per rank, partial results sent back by each rank, etc.)

- (2) A pdf file using the following name: lastNamefirstNameLab1.pdf

- Any information that you want to give me (e.g., what worked, what didn't, etc.)
- **Part 1:** What the master rank printed (at the end): wcalcTime for each worker rank, the maximum wcalcTime, the sum of the wcalcTime for all worker ranks, and the total time_spent (time from sending the data to workers until the master have all the results).
- **Part 2:** Two tables
 - Table 1 (for N = 100000) - Row labels: R (# of ranks) = 4, 8, 16, 32; Column labels: estimated result, time taken
 - Table 2 (for R = 32) - Row labels: N = 100, 1000, ... , 100000 ; Column labels: estimated result, time taken
- **Part 3:** Two tables (same as Part 2)